

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("Fitness.csv")
# .dropna(axis="columns")
df
```

```
Out[2]:
```

	Row Labels	Sum of Jan	Sum of Feb	Sum of Mar	Sum of Total Sales
0	A	5.62%	7.73%	6.16%	75
1	B	4.21%	17.27%	19.21%	160
2	C	9.83%	11.60%	5.17%	101
3	D	2.81%	21.91%	7.88%	127
4	E	25.28%	10.57%	11.82%	179
5	F	8.15%	16.24%	18.47%	167
6	G	18.54%	8.76%	17.49%	171
7	H	25.56%	5.93%	13.79%	170
8	Grand Total	100.00%	100.00%	100.00%	1150

```
In [3]: df.head()
```

```
Out[3]:
```

	Row Labels	Sum of Jan	Sum of Feb	Sum of Mar	Sum of Total Sales
0	A	5.62%	7.73%	6.16%	75
1	B	4.21%	17.27%	19.21%	160
2	C	9.83%	11.60%	5.17%	101
3	D	2.81%	21.91%	7.88%	127
4	E	25.28%	10.57%	11.82%	179

## Data cleaning and pre processing

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9 entries, 0 to 8
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Row Labels            9 non-null     object
1   Sum of Jan            9 non-null     object
2   Sum of Feb            9 non-null     object
3   Sum of Mar            9 non-null     object
4   Sum of Total Sales    9 non-null     int64
dtypes: int64(1), object(4)
memory usage: 488.0+ bytes
```

In [5]: `df.describe()`

Out[5]:

	Sum of Total Sales
count	9.000000
mean	255.555556
std	337.332963
min	75.000000
25%	127.000000
50%	167.000000
75%	171.000000
max	1150.000000


In [6]: `df.columns`

Out[6]: Index(['Row Labels', 'Sum of Jan', 'Sum of Feb', 'Sum of Mar',  
'Sum of Total Sales'],  
dtype='object')

## EDA and VISUALIZATION

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1de304954c0>
```




```
In [8]: sns.distplot(df["Sum of Total Sales"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
Out[8]: <AxesSubplot:xlabel='Sum of Total Sales', ylabel='Density'>
```



```
In [9]: df1 = df[['Row Labels', 'Sum of Jan', 'Sum of Feb', 'Sum of Mar',  
                'Sum of Total Sales']]
```

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x = df1[['Sum of Total Sales', 'Sum of Total Sales']]
y = df1['Sum of Total Sales']
```

## split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

```
Out[13]: LinearRegression()
```

```
In [14]: lr.intercept_
```

```
Out[14]: 0.0
```


```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns = ['Co-efficient'])
coeff
```

```
Out[15]:
```

	Co-efficient
Sum of Total Sales	0.5
Sum of Total Sales	0.5

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x1de30eeb4f0>



```
In [17]: lr.score(x_test,y_test)
```

Out[17]: 1.0

## ACURACY

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

Out[19]: 0.9999999999647301

```
In [20]: rr.score(x_test,y_test)
```

Out[20]: 0.9999999994038117

```
In [21]: la = Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[21]: Lasso(alpha=10)

```
In [22]: la.score(x_test,y_test)
```

Out[22]: 0.9999999141478655

```
In [23]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

Out[23]: ElasticNet()

```
In [24]: print(en.coef_)

[9.99989310e-01 7.12655515e-06]
```

```
In [25]: print(en.intercept_)

0.0011206747564642683
```

```
In [26]: print(en.predict(x_test))

[167.00052559 171.00051134 75.00085342]
```

```
In [27]: print(en.score(x_test,y_test))

0.9999999997853666
```

```
In [28]: # Evaluation Metrics
from sklearn import metrics
```

```
In [29]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))

Mean Absolute Error: 0.0
```

```
In [30]: print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))

Root Mean Squared Error: 0.0
```