

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("Horse.csv")
# .dropna(axis="columns")
df
```

Out[2]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	Coun
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	Sver
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	Sver
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	Sver
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	Sver
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	Sver
...
27003	14.06.2020	Sha Tin	11	1200	Gress	1450000	6	A Hamelin	59	Austr
27004	21.06.2020	Sha Tin	2	1200	Gress	967000	7	K C Leung	57	Austr
27005	21.06.2020	Sha Tin	4	1200	Gress	967000	6	Blake Shinn	57	Austr
27006	21.06.2020	Sha Tin	5	1200	Gress	967000	14	Joao Moreira	57	N Zeala
27007	21.06.2020	Sha Tin	11	1200	Gress	1450000	7	C Schofield	55	N Zeala

27008 rows × 21 columns



In [3]: `df.head()`

Out[3]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	Country	..
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	Sverige	..
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	Sverige	..
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	Sverige	..
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	Sverige	..
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	Sverige	..

5 rows × 21 columns



Data cleaning and pre processing

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27008 entries, 0 to 27007
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Dato                   27008 non-null  object
1   Track                 27008 non-null  object
2   Race Number           27008 non-null  int64
3   Distance              27008 non-null  int64
4   Surface               27008 non-null  object
5   Prize money           27008 non-null  int64
6   Starting position     27008 non-null  int64
7   Jockey                27008 non-null  object
8   Jockey weight         27008 non-null  int64
9   Country               27008 non-null  object
10  Horse age             27008 non-null  int64
11  TrainerName           27008 non-null  object
12  Race time             27008 non-null  object
13  Path                  27008 non-null  int64
14  Final place           27008 non-null  int64
15  FGrating              27008 non-null  int64
16  Odds                  27008 non-null  object
17  RaceType              27008 non-null  object
18  HorseId               27008 non-null  int64
19  JockeyId              27008 non-null  int64
20  TrainerID             27008 non-null  int64
dtypes: int64(12), object(9)
memory usage: 4.3+ MB
```

In [5]: `df.describe()`

Out[5]:

	Race Number	Distance	Prize money	Starting position	Jockey weight	Horse age
count	27008.000000	27008.000000	2.700800e+04	27008.000000	27008.000000	27008.000000
mean	5.268624	1401.666173	1.479445e+06	6.741447	55.867373	5.246408
std	2.780088	276.065045	2.162109e+06	3.691071	2.737006	1.519880
min	1.000000	1000.000000	6.600000e+05	1.000000	47.000000	2.000000
25%	3.000000	1200.000000	9.200000e+05	4.000000	54.000000	4.000000
50%	5.000000	1400.000000	9.670000e+05	7.000000	56.000000	5.000000
75%	8.000000	1650.000000	1.450000e+06	10.000000	58.000000	6.000000
max	11.000000	2400.000000	2.800000e+07	14.000000	63.000000	12.000000

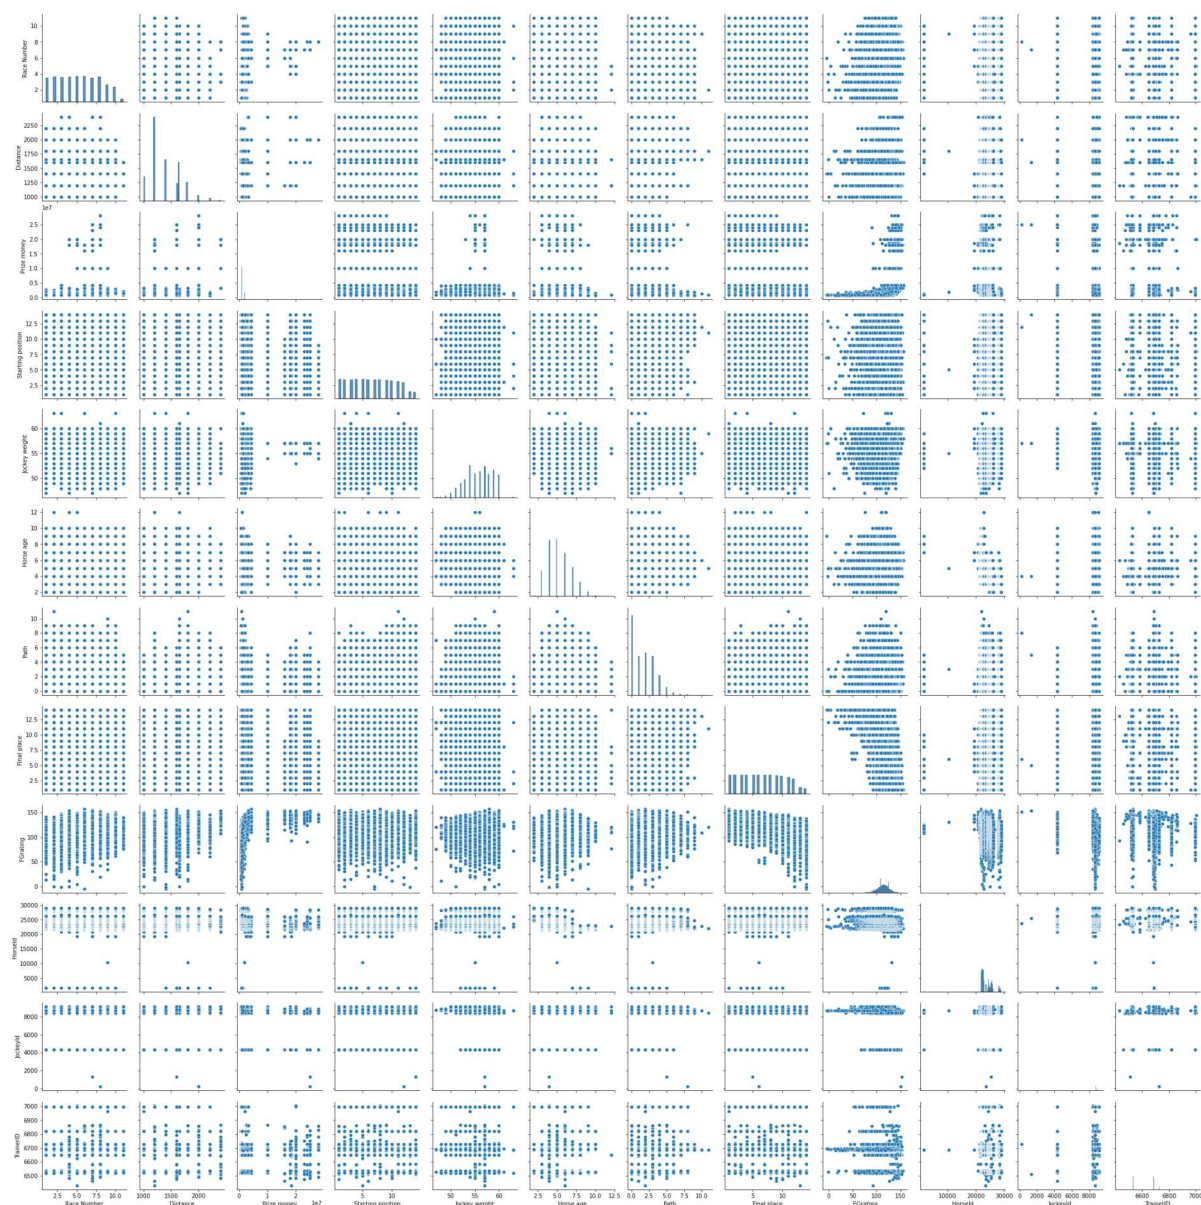
In [6]: `df.columns`

Out[6]: Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',
'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',
'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',
'RaceType', 'HorseId', 'JockeyId', 'TrainerID'],
dtype='object')

EDA and VISUALIZATION

```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x16381cb8b50>
```




```
In [8]: sns.distplot(df['Distance'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```


```
Out[8]: <AxesSubplot:xlabel='Distance', ylabel='Density'>
```



```
In [9]: df1 = df[['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',
                  'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',
                  'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',
                  'RaceType', 'HorseId', 'JockeyId', 'TrainerID']]
```

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: x = df1[['Race Number', 'Distance', 'Prize money', 'Starting position', 'Jockey weight']]
y = df1['FGRating']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

Out[13]: LinearRegression()

```
In [14]: lr.intercept_
```

Out[14]: 101.38500963011558


```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns = ['Co-efficient'])
coeff
```

Out[15]:

	Co-efficient
Race Number	1.997801
Distance	0.000797
Prize money	0.000001
Starting position	0.014436
Jockey weight	1.048928
Horse age	0.134051
Path	1.247377
Final place	-1.985124
Horseld	-0.000288
JockeyId	-0.000060
TrainerID	-0.006380

```
In [16]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x16391246eb0>



```
In [17]: lr.score(x_test,y_test)
```

Out[17]: 0.6438692690608674

ACURACY

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

Out[19]: 0.6411267182253932

```
In [20]: rr.score(x_test,y_test)
```

Out[20]: 0.6438691554560037

```
In [21]: la = Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[21]: Lasso(alpha=10)

```
In [22]: la.score(x_test,y_test)
```

Out[22]: 0.44887089282073434

```
In [23]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

Out[23]: ElasticNet()

```
In [24]: print(en.coef_)
```

```
[ 1.79018006e+00  1.52806735e-03  1.47647609e-06  2.41437525e-02
 9.16744500e-01  0.00000000e+00  8.50859987e-01 -1.86454844e+00
-3.53067249e-04 -1.00839978e-04 -6.48166348e-03]
```

```
In [25]: print(en.intercept_)
```

111.8594183134557

```
In [26]: print(en.predict(x_test))
```

```
[129.42015878 115.09866688 112.13041113 ... 100.25105845  99.70673682
124.60852832]
```

```
In [27]: print(en.score(x_test,y_test))
```

0.6380379502227375

```
In [28]: # Evaluation Metrics
from sklearn import metrics
```

```
In [29]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 5.5874718729549135

```
In [30]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 63.93305087761121

```
In [31]: print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 7.995814584994528