

```
In [8]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [9]: df = pd.read_csv("world.csv").dropna()

df
```

Out[9]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Calling Code	Capital/M
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	93.0	K
1	Albania	105	AL	43.10%	28,748	9,000	11.78	355.0	Ti
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	213.0	Al
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	244.0	Lu
6	Argentina	17	AR	54.30%	2,780,400	105,000	17.02	54.0	Buenos /
...	...	...	...	...	...	...	...	...	...
185	United Kingdom	281	GB	71.70%	243,610	148,000	11.00	44.0	Loi
186	United States	36	US	44.40%	9,833,517	1,359,000	11.60	1.0	Washin
187	Uruguay	20	UY	82.60%	176,215	22,000	13.86	598.0	Montev
191	Vietnam	314	VN	39.30%	331,210	522,000	16.75	84.0	H
193	Zambia	25	ZM	32.10%	752,618	16,000	36.19	260.0	Lu

110 rows × 35 columns



```
In [10]: df.head()
```

Out[10]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land( %)	Land Area(Km2)	Armed Forces size	Birth Rate	Calling Code	Capital/Major City
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	93.0	Kabu
1	Albania	105	AL	43.10%	28,748	9,000	11.78	355.0	Tirana
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	213.0	Algiers
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	244.0	Luanda
6	Argentina	17	AR	54.30%	2,780,400	105,000	17.02	54.0	Buenos Aires

5 rows × 35 columns



## Data cleaning and pre processing

In [11]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 110 entries, 0 to 193
Data columns (total 35 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Country                                                                110 non-null   object
1   Density                                                                110 non-null   object
   (P/Km2)
2   Abbreviation                                                            110 non-null   object
3   Agricultural Land( %)                                                  110 non-null   object
4   Land Area(Km2)                                                         110 non-null   object
5   Armed Forces size                                                      110 non-null   object
6   Birth Rate                                                             110 non-null   float64
7   Calling Code                                                           110 non-null   float64
8   Capital/Major City                                                    110 non-null   object
9   Co2-Emissions                                                         110 non-null   object
10  CPI                                                                    110 non-null   object
11  CPI Change (%)                                                         110 non-null   object
12  Currency-Code                                                         110 non-null   object
13  Fertility Rate                                                         110 non-null   float64
14  Forested Area (%)                                                      110 non-null   object
15  Gasoline Price                                                         110 non-null   object
16  GDP                                                                    110 non-null   object
17  Gross primary education enrollment (%) 110 non-null   object
18  Gross tertiary education enrollment (%) 110 non-null   object
19  Infant mortality                                                       110 non-null   float64
20  Largest city                                                           110 non-null   object
21  Life expectancy                                                        110 non-null   float64
22  Maternal mortality ratio                                              110 non-null   float64
23  Minimum wage                                                           110 non-null   object
24  Official language                                                      110 non-null   object
25  Out of pocket health expenditure  110 non-null   object
26  Physicians per thousand                                               110 non-null   float64
27  Population                                                             110 non-null   object
28  Population: Labor force participation (%) 110 non-null   object
29  Tax revenue (%)                                                        110 non-null   object
30  Total tax rate                                                         110 non-null   object
31  Unemployment rate                                                     110 non-null   object
32  Urban_population                                                      110 non-null   object
33  Latitude                                                               110 non-null   float64
34  Longitude                                                             110 non-null   float64
dtypes: float64(9), object(26)
memory usage: 30.9+ KB
```

In [12]: `df.describe()`

Out[12]:

	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand	Latitude	
<b>count</b>	110.000000	110.000000	110.000000	110.000000	110.000000	110.000000	110.000000	110.000000	1
<b>mean</b>	20.196455	344.290909	2.672182	20.271818	72.671818	137.227273	1.919182	20.362677	
<b>std</b>	10.039056	341.231562	1.308142	18.453214	7.000788	201.171462	1.598116	24.432140	
<b>min</b>	6.400000	1.000000	0.980000	1.700000	54.300000	2.000000	0.010000	-40.900557	-1
<b>25%</b>	11.075000	70.000000	1.682500	6.100000	67.625000	15.250000	0.467500	7.623255	
<b>50%</b>	17.830000	239.500000	2.200000	13.600000	74.400000	41.000000	1.640000	21.033608	
<b>75%</b>	27.962500	420.750000	3.505000	31.500000	77.350000	176.000000	3.007500	40.124603	
<b>max</b>	46.080000	1876.000000	6.910000	78.500000	83.300000	1120.000000	7.120000	61.524010	1

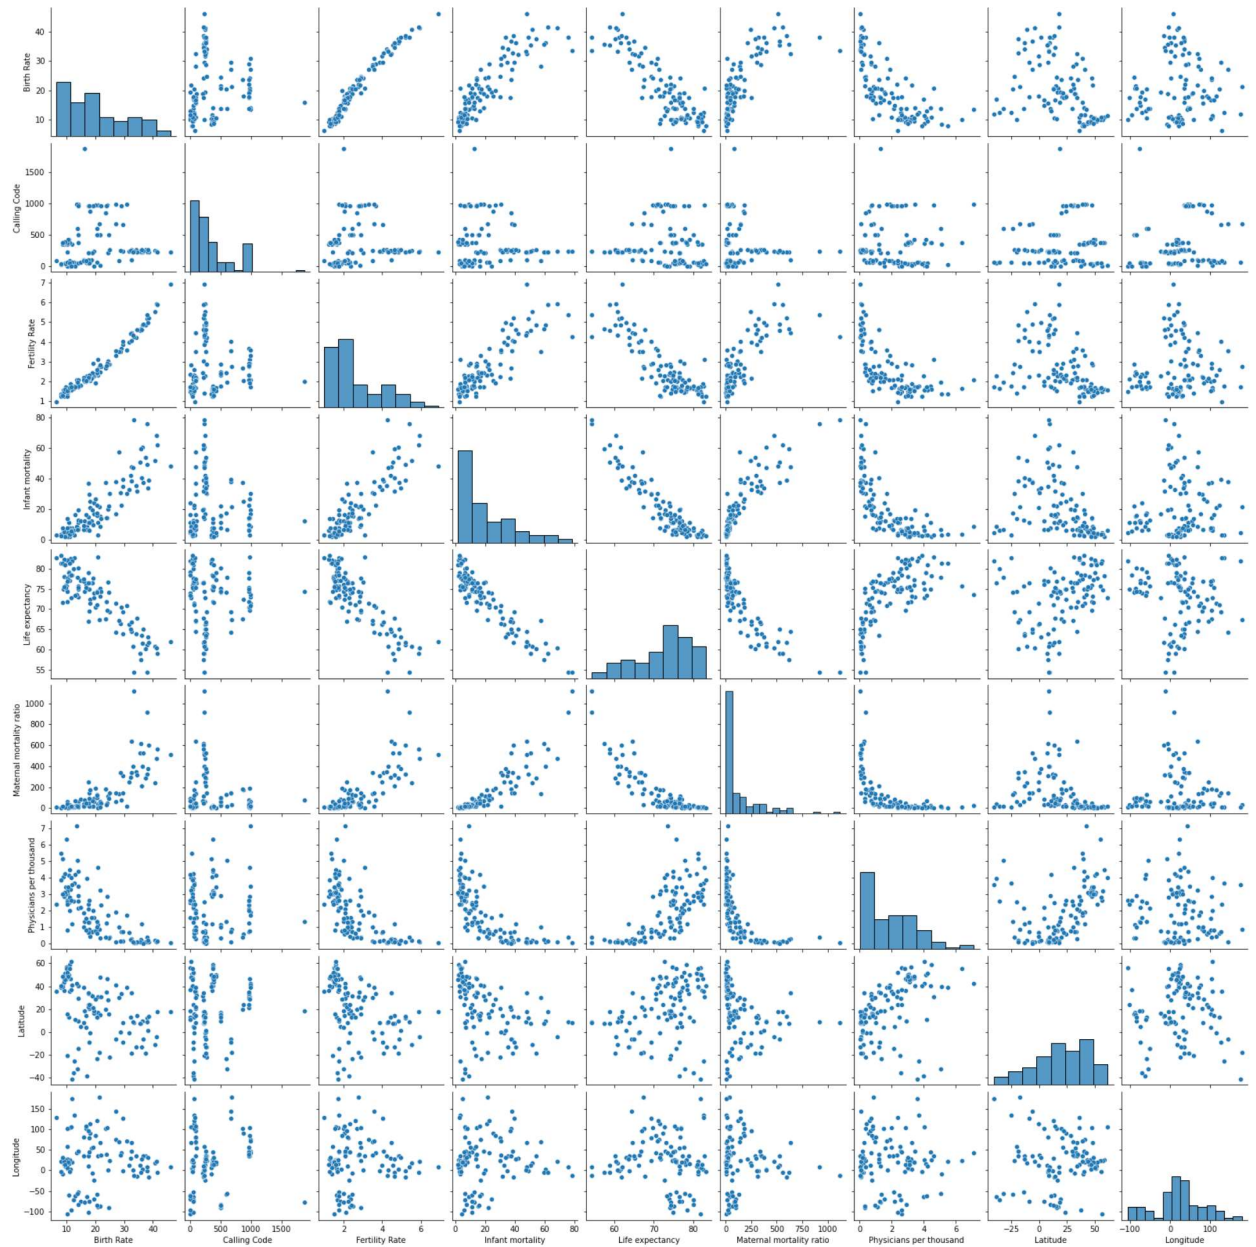
In [13]: `df.columns`

Out[13]: Index(['Country', 'Density\n(P/Km2)', 'Abbreviation', 'Agricultural Land( %)', 'Land Area(Km2)', 'Armed Forces size', 'Birth Rate', 'Calling Code', 'Capital/Major City', 'Co2-Emissions', 'CPI', 'CPI Change (%)', 'Currency-Code', 'Fertility Rate', 'Forested Area (%)', 'Gasoline Price', 'GDP', 'Gross primary education enrollment (%)', 'Gross tertiary education enrollment (%)', 'Infant mortality', 'Largest city', 'Life expectancy', 'Maternal mortality ratio', 'Minimum wage', 'Official language', 'Out of pocket health expenditure', 'Physicians per thousand', 'Population', 'Population: Labor force participation (%)', 'Tax revenue (%)', 'Total tax rate', 'Unemployment rate', 'Urban\_population', 'Latitude', 'Longitude'], dtype='object')

## EDA and VISUALIZATION

```
In [14]: sns.pairplot(df)
```


```
Out[14]: <seaborn.axisgrid.PairGrid at 0x206177becd0>
```



```
In [15]: sns.distplot(df['Latitude'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

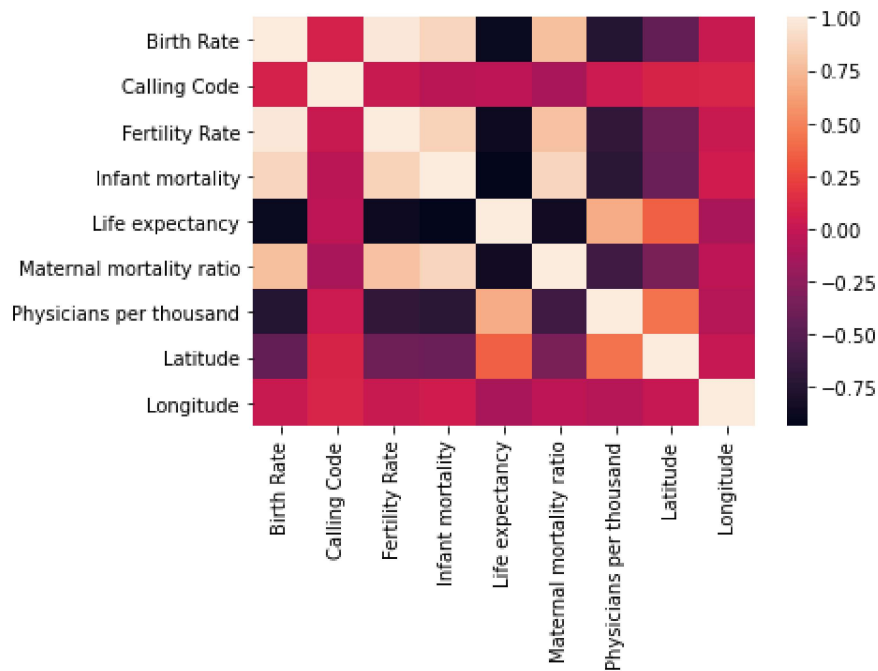
```
Out[15]: <AxesSubplot:xlabel='Latitude', ylabel='Density'>
```



```
In [16]: df1 = df[['Country', 'Density\n(P/Km2)', 'Abbreviation', 'Agricultural Land( %)',
'Land Area(Km2)', 'Armed Forces size', 'Birth Rate', 'Calling Code',
'Capital/Major City', 'Co2-Emissions', 'CPI', 'CPI Change (%)',
'Currency-Code', 'Fertility Rate', 'Forested Area (%)',
'Gasoline Price', 'GDP', 'Gross primary education enrollment (%)',
'Gross tertiary education enrollment (%)', 'Infant mortality',
'Largest city', 'Life expectancy', 'Maternal mortality ratio',
'Minimum wage', 'Official language', 'Out of pocket health expenditure',
'Physicians per thousand', 'Population',
'Population: Labor force participation (%)', 'Tax revenue (%)',
'Total tax rate', 'Unemployment rate', 'Urban_population', 'Latitude',
'Longitude']]
```

```
In [17]: sns.heatmap(df1.corr())
```

```
Out[17]: <AxesSubplot:>
```



```
In [18]: x = df1[['Birth Rate','Calling Code','Fertility Rate','Infant mortality','Life expectancy',
y = df1['Physicians per thousand']
```

## split the data into training and test data

```
In [19]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [20]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

```
Out[20]: LinearRegression()
```

```
In [21]: lr.intercept_
```

```
Out[21]: 1.1102230246251565e-14
```


```
In [22]: coeff = pd.DataFrame(lr.coef_, x.columns, columns = ['Fertility Rate'])
coeff
```

Out[22]:

	Fertility Rate
Birth Rate	1.554247e-16
Calling Code	-1.763489e-18
Fertility Rate	1.198357e-15
Infant mortality	-2.629640e-16
Life expectancy	-1.627314e-16
Maternal mortality ratio	5.771328e-18
Physicians per thousand	1.000000e+00
Latitude	-7.941307e-18
Longitude	-4.445337e-18

```
In [23]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[23]: <matplotlib.collections.PathCollection at 0x2061bf1d730>



```
In [24]: lr.score(x_test,y_test)
```

Out[24]: 1.0

## ACURACY

```
In [25]: from sklearn.linear_model import Ridge,Lasso
```

```
In [26]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

Out[26]: 0.994931296508725

```
In [27]: rr.score(x_test,y_test)
```

```
Out[27]: 0.9948764479359316
```

```
In [28]: la = Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[28]: Lasso(alpha=10)
```

```
In [29]: la.score(x_test,y_test)
```

```
Out[29]: 0.46994638257697163
```

```
In [30]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

```
Out[30]: ElasticNet()
```

```
In [31]: print(en.intercept_)
```

```
1.944500190641182
```

```
In [32]: print(en.predict(x_test))
```

```
[ 3.53045825  2.68122952  2.39110232  1.81993614  3.07893568  2.41689141
 1.33934525  0.2978623   2.92878239  3.24377094 -0.16892324  1.38545523
 1.72878448  3.11067205  0.09913052  0.46238285  2.62237527  3.2256435
 1.78924011  3.389855   3.23199662  1.1803987   2.60418248  0.2178042
 2.95030164  0.06101522  2.2288558   3.22648357  1.23939681  0.19362643
 3.60575579  2.2116378   0.45649973]
```

```
In [33]: print(en.score(x_test,y_test))
```

```
0.8396666269015641
```

```
In [34]: # Evaluation Metrics
from sklearn import metrics
```

```
In [35]: print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error: 1.5444295673999573e-15
```

```
In [36]: print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error: 4.991031401463319e-30
```

```
In [38]: print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
Root Mean Squared Error: 2.234061637794114e-15
```



